

Edge AI based Federated Learning Architecture for Healthcare Applications

A Project Report Submitted
in Partial Fulfilment of the Requirements
for the Degree of

BACHERLORS OF TECHNOLOGY
in
COMPUTER SCIENCE AND ENGINEERING

by

HARSHITHA DONDETI (2022BCS0209)
MADHAV MURALI (2022BCS0050)
MADDELA MANOJ VENKATA SAI (2022BCS0126)
THADI SUSHMITHA MANI (2022BCD005)



to

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
INDIAN INSTITUTE OF INFORMATION TECHNOLOGY
KOTTAYAM-686635, INDIA

November 2025

DECLARATION

We, **HARSHITHA DONDETI (2022BCS0209)**, **MADHAV MURALI (2022BCS0050)**, **MADDELA MANOJ VENKATA SAI (2022BCS0126)** and **THADI SUSHMITHA MANI (2022BCD005)**, hereby declares that this report entitled “**Edge AI based Federated Learning Architecture for Healthcare Applications**” submitted to Indian Institute of Information Technology Kottayam towards partial requirement of **Bachelor of Technology in COMPUTER SCIENCE AND ENGINEERING** is an original work carried out by me under the supervision of **Dr. Koppala Guravaiah** and has not formed the basis for the award of any degree or diploma, in this or any other institution or university. I have sincerely tried to uphold academic ethics and honesty. Whenever an external information, statement, or result is used, that have been duly acknowledged and cited.

Kottayam-686635

November 2025

HARSHITHA DONDETI (2022BCS0209)

MADHAV MURALI (2022BCS0050)

MADDELA MANOJ VENKATA SAI (2022BCS0126)

THADI SUSHMITHA MANI (2022BCS0005)

CERTIFICATE

This is to certify that the work contained in this project report entitled “**Edge AI based Federated Learning Architecture for Healthcare Applications**” submitted by (**HARSHITHA DONDETI (2022BCS0209)**, **MADHAV MURALI (2022BCS0050)**, **MADDELA MANOJ VENKATA SAI (2022BCS0126)** and **THADI SUSHMITHA MANI (2022BCD005)**) to the Indian Institute of Information Technology Kottayam towards the partial requirement of **Bachelor of Technology** in **INDIAN INSTITUTE OF INFORMATION TECHNOLOGY KOTTAYAM** has been carried out by them under my supervision and that it has not been submitted elsewhere for the award of any degree.

Kottayam-686635

November 2025

(Dr. Koppala Guravaiah)

Project Supervisor

ABSTRACT

The recent use of artificial intelligence in the medical field has transformed the monitoring and diagnostics of the patient and has also generated significant issues concerning data privacy, security and scale. The conventional cloud-based healthcare AI systems rely on data being stored centrally, which brings additional communication overhead and latency to the system and makes it susceptible to breach of data. To address these constraints, this project suggests a Federated Learning Architecture based on Edge AI and Healthcare Applications that will allow training a model collaboratively in multiple hospitals without disclosing sensitive patient data. The system proposed carries out local training on devices of edges with the use of electronic health records (EHRs) and medical imaging data and only conveys model weights to a worldwide server to be aggregated. This method guarantees that there is confidentiality of data, and high muscle of the model. To enhance the interoperability and scalability of heterogeneous healthcare systems, the architecture includes lightweight communication protocols (HTTPS, MQTTS, BLE) and FHIR standards. The system will seek to attain privacy-sustaining, resource-effective, and agile healthcare intelligence that is apt in both urban and rural medical systems.

Contents

List of Figures	ix
List of Tables	x
1 Introduction	1
1.1 Existing Healthcare Infrastructure	1
1.2 Limitations in the Present Approach	2
1.3 Proposed Decentralized Intelligent Framework	2
1.4 Integration of Digital Twin Technology	3
2 Literature Review	4
3 Methodology	6
3.1 System Architecture	6
3.1.1 Core Components	6
3.2 Workflow Process	7
3.2.1 Data Acquisition and Digital Twin Creation	7
3.2.2 Local Training on Edge Devices	8
3.2.3 Model Comparison and Deployment	8

3.2.4	Federated Learning Cycle	8
3.2.5	Continuous Feedback Loop	8
3.3	Data Processing Pipeline	9
3.3.1	Datasets Used	9
3.3.2	Data Preprocessing	9
3.4	Model Architectures	10
3.4.1	CNN-LSTM for Chest X-ray Classification	10
3.4.2	CNN-BiLSTM for ECG Classification	10
3.5	Federated Learning Implementation	11
3.5.1	Configuration Parameters	12
3.6	Evaluation Framework	13
4	Experimental Setup and Implementation	14
4.1	Datasets and Preprocessing	14
4.1.1	NIH Chest X-ray14 Dataset	14
4.1.2	MIT-BIH Arrhythmia Database	14
4.2	Data Preprocessing Pipeline	16
4.2.1	Chest X-ray Preprocessing	16
4.2.2	ECG Signal Preprocessing	17
4.3	Non-IID Data Simulation	17
4.4	Federated Learning Configuration	18
4.4.1	Training Parameters	18
4.5	Implementation Details	18
4.5.1	Software Stack	18
4.5.2	Development Environment	19

5	Results and Discussion	20
5.1	Chest X-ray Classification Results	20
5.1.1	Overall Performance	20
5.1.2	Client Performance Analysis	21
5.2	ECG Classification Results	22
5.2.1	Heartbeat Classification Performance	22
5.2.2	Abnormal Beat Detection	22
5.3	Comparison with Baseline Approaches	23
5.3.1	Performance Comparison	23
5.4	Key Findings and Insights	24
5.4.1	Main Contributions	24
5.4.2	Limitations and Challenges	24
6	Conclusion and Future Work	26
6.1	Summary of Contributions	26
6.1.1	Technical Innovations and Performance	26
6.2	Key Findings	27
6.2.1	Privacy-Performance Balance	27
6.2.2	Data Heterogeneity Management	28
6.3	Limitations and Challenges	28
6.3.1	Technical Limitations	28
6.3.2	Practical Challenges	29
6.4	Future Research Directions	29
6.4.1	Immediate Enhancements for the near future:	29
6.5	Concluding Remarks	30

List of Figures

3.1	Proposed Federated Learning System Architecture	7
4.1	Sample Chest X-ray Images from NIH Dataset	16
4.2	ECG Waveform Examples from MIT-BIH Database	16
5.1	Convergence of Global Model for Chest X-ray Classification . .	21

List of Tables

2.1	Summary of Federated Learning Literature in Healthcare . . .	5
3.1	Medical Datasets for Validation	9
3.2	CNN-LSTM Architecture for X-ray Classification	10
3.3	CNN-BiLSTM Architecture for ECG Classification	11
3.4	Federated Learning Parameters	12
3.5	Evaluation Metrics	13
4.1	NIH Chest X-ray14 Dataset Composition	15
4.2	MIT-BIH Dataset Beat Distribution	15
4.3	Non-IID Data Distribution Strategy	17
4.4	Training Configuration Parameters	18
4.5	Software Components and Versions	19
5.1	Federated Learning Performance on Chest X-ray Dataset . . .	20
5.2	Individual Client Performance on Chest X-ray Data	21
5.3	ECG Heartbeat Classification Results	22
5.4	Abnormal Heartbeat Detection Performance	23
5.5	Comparison with Alternative Approaches	23

6.1	Key Technical Contributions and Their Impact	27
6.2	Performance Summary Across Medical Domains	27
6.3	Identified Limitations and Mitigation Strategies	28
6.4	Near-term Research Priorities	29

Chapter 1

Introduction

Automated diagnostics, real-time monitoring, and data-driven decisions are considered to be the revolution in healthcare driven by rapid development of artificial intelligence (AI) and Internet of Medical Things (IoMT). Techniques based on smart systems are now employed to process electronic health records (EHRs), medical images as well as sensor data in hospitals. Although the technologies increase efficiency, there are some important issues of critical concern associated with data privacy, latency, and interoperability. The use of cloud-based systems also reveals serious flaws emphasizing efficient, decentralized, and secure solutions in healthcare.

1.1 Existing Healthcare Infrastructure

The existing healthcare networks rely much on artificial intelligence (AI) application platforms held on clouds interrelated with smart medical devices of wearables and sensors. Such devices constantly record patient information

and send them to the centralized cloud servers to process them. Even though the model is helpful, it has too much network dependency and centralization. Consequently, it repeatedly lacks support of real-time responses and makes a considerable load on the bandwidth and cloud storage services.

1.2 Limitations in the Present Approach

There are several challenges associated with centralized healthcare architectures which limit their effectiveness as well as scalability. The first priority is privacy and safety of data, storing sensitive patient records in central servers puts more risk on the breach of data. Also, due to continuous data transmission, the system is unsuitable to emergency or real-time medical decision-making because of high latency and bandwidth usage. In addition, the interoperability and standardization problems are caused by compliance with various privacy laws, and unequal electronic health record (EHR) designs. The lack of homogenous data on hospitals also undermines the efficacy of traditional AI models because they are trained on unequal and unbalanced data.

1.3 Proposed Decentralized Intelligent Framework

In order to eliminate these constraints, this project presents an Edge AI-backed Federated Learning framework on healthcare applications. The proposed system will allow hospitals and medical devices to work with their

own data to train AI models locally instead of centralizing data. The trained model parameters are the only ones that are shared with a central aggregator, to update model around the world, which guarantees full data privacy. This decentralization will result in a hugely lower load on the network, and communication costs, and high-quality models. Its implementation employs lightweight but secure communication schemes which render it applicable to the implementation in bandwidth constrained and rural healthcare settings. Moreover, with FHIR (Fast Healthcare Interoperability Resources), it should be seen that it is standardized and can be interoperable with various medical systems.

1.4 Integration of Digital Twin Technology

The main new element of the suggested architecture is the inclusion of Digital Twin technology, which develops a virtual replica of a patient or a healthcare process. These digital twins constantly reflect a real situation of the patient ,helps monitor them in real time, provide predictive analysis, and tailor treatment. The connection between Digital Twins and Edge AI and Federated Learning could means with its help, proactive healthcare management, a smaller time of diagnosis, and assisting in decision making is possible due to the use of simulated clinical scenarios.

Chapter 2

Literature Review

Federated Learning (FL) has emerged as a transformative approach for health-care AI, enabling collaborative model training across institutions while keeping patient data local. This addresses critical privacy concerns in traditional centralized systems. Current research demonstrates FL's effectiveness across medical imaging, electronic health records, and IoT monitoring, though challenges remain in handling data heterogeneity, privacy-accuracy tradeoffs, and real-world deployment scalability.

The literature indicates FL's strong potential for privacy-preserving health-care AI, though key challenges around data heterogeneity, communication efficiency, and clinical integration require further research. Future work should focus on adaptive aggregation strategies and real-world validation to bridge the gap between theoretical promise and practical implementation.

Reference	Main Contribution	Key Challenges	Future Directions
Ng et al. [9]	FL for hospitals with small datasets	Model aggregation issues, data format variations	Better aggregation methods, improved interoperability
Nguyen et al. [10]	Comprehensive FL survey for healthcare	Non-IID data, communication costs, device variety	Lightweight models, energy-efficient frameworks
Antunes et al. [2]	EHR-based FL architecture	Data heterogeneity, privacy-accuracy balance	Differential privacy, hardware security features
Rahman et al. [14]	FL application taxonomy	Multi-institution privacy, scaling limitations	Blockchain integration, explainable AI systems
Rauniyar et al. [15]	Medical FL taxonomy	Communication costs, privacy risks	Lightweight FL, multimodal learning approaches
Bashir et al. [3]	FL in healthcare meta-verse	Multimodal data, governance issues	Energy-efficient, interoperable systems
Reddy et al. [16]	FL frameworks survey	Privacy-accuracy trade-offs, trust issues	Transfer learning, benchmark datasets
Elayan et al. [5]	IoT healthcare with FL	Biased data, slow convergence, limited resources	Adaptive feedback, data balance methods
Polap et al. [13]	FL with blockchain	High communication needs, latency problems	Scalability improvements, reduced complexity
Pace et al. [11]	Edge computing framework	Missing AI features, scaling problems	Edge-FL combination
Dayan et al. [4]	COVID-19 prediction across hospitals	Data biases, interoperability issues	Clinical testing, fair evaluation methods
Patel et al. [12]	FL adoption analysis	Limited deployment, regulatory compliance	Standard frameworks, optimized protocols
Alazab et al. [1]	Digital twin review	Compatibility gaps, energy waste	AI-enhanced twins, secure IoT systems
Stelter et al. [17]	Privacy impact on FL	Accuracy reduction from privacy methods	Synthetic data generation
Hernandez-Cruz et al. [7]	FL for medical images	Non-IID data, limited datasets	Practical solutions, diverse datasets
Kapila et al. [8]	FL with feature selection	Simulation-only testing, lacks real use	Explainable AI, privacy methods
Haripriya et al. [6]	Privacy-preserving FL for images	Communication overhead, scaling limits	Multimodal FL, efficient aggregation

Table 2.1: Summary of Federated⁵ Learning Literature in Healthcare

Chapter 3

Methodology

3.1 System Architecture

We designed a comprehensive Edge AI-based Federated Learning system that addresses the limitations of traditional cloud-based healthcare AI. The architecture enables collaborative learning across hospitals while keeping patient data private and local.

3.1.1 Core Components

Edge Nodes: These would be the individual hospitals or medical devices that store and process local patient data. Then train AI models on-premises. They are the ones supposed to maintain data privacy and security. Should Support multiple data types like ECG, X-ray, EHR etc..

Aggregation Server: Central coordinator that distributes initial global models and aggregates model weights using FedAvg or other federated learning methods. They also manages training rounds and convergence. But they




Figure 3.1: Proposed Federated Learning System Architecture

should never access raw patient data

Communication Layer: Secure protocols including HTTPS with TLS 1.3 for parameter transmission, MQTTS for lightweight messaging, FHIR standards for healthcare data interoperability are to be used.

3.2 Workflow Process

Our federated learning workflow consists of the following five key steps:

3.2.1 Data Acquisition and Digital Twin Creation

Collecting EHR data from hospital systems. Then preprocess and clean medical data. Create digital twin representations of patients/systems followed by converting raw data into standardized formats

3.2.2 Local Training on Edge Devices

Train child models on local hospital data. Use labeled data for supervised learning. These are then used for detecting and correcting mislabeled data using inference errors. Optimize models for local data characteristics.

3.2.3 Model Comparison and Deployment

We compare child model performance with global model and then deploy better performing model locally. But we do maintain global model as fallback option to ensure continuous service availability

3.2.4 Federated Learning Cycle

Share model weights (not data) with aggregation server and update global model using Federated Averaging (FedAvg). Redistribute improved global model to all nodes and use transfer learning to enhance local models.

3.2.5 Continuous Feedback Loop

Reuse inference results for model improvement. Identify and correct mislabeled data. Use shadow models for experimentation. Maintain continuous learning and adaptation

3.3 Data Processing Pipeline

3.3.1 Datasets Used

We validated our approach using two major medical datasets:

Table 3.1: Medical Datasets for Validation

Dataset			Type	Size	Application
NIH Chest X-ray14	X-		Medical images	112,120 images	14 disease classification
MIT-BIH Arrhythmia	Ar-		ECG signals	10,729 beats	Heartbeat classification

3.3.2 Data Preprocessing

For NIH Chest X-ray14: Downloaded via KaggleHub API. First filter valid PNG files were resized to 224×224 pixels and then normalize pixel values using image-net statistics.

For MIT-BIH Arrhythmia: Downloaded records 100-104 using WFDB. Removed noise and artifacts to discard invalid beats. Used Detrend and z-score normalize signals. Segmented into 360-sample windows around R-peaks.

Table 3.2: CNN-LSTM Architecture for X-ray Classification

Layer	Parameters	Output Shape	Purpose
Input	-	(224, 224, 3)	Raw X-ray image
Conv2D + Pooling	32 filters, 3×3	(112, 112, 32)	Feature extraction
Conv2D + Pooling	64 filters, 3×3	(56, 56, 64)	Spatial patterns
Flatten	-	(200704)	Prepare for LSTM
Reshape	-	(1, 200704)	Sequence formation
LSTM	128 units	(128)	Temporal dependencies
Dense	64 units	(64)	Feature combination
Output	14 units	(14)	Multi-label classification

3.4 Model Architectures

3.4.1 CNN-LSTM for Chest X-ray Classification

Key Features: We used Sigmoid activation for multi-label classification. Binary Cross Entropy loss function and Adam optimizer (learning rate 1e-4) with a Batch size: 32.

3.4.2 CNN-BiLSTM for ECG Classification

Key Features: Used Focal Loss to handle class imbalance. Adam optimizer with learning rate 1e-3. Attention mechanism for important time-steps. Dropout rates: 0.3 (CNN), 0.4 (FC) for reducing overfitting.

Table 3.3: CNN-BiLSTM Architecture for ECG Classification

Layer	Parameters	Output Shape	Purpose
Input	-	(360, 1)	ECG signal segment
Conv1D	32 filters, kernel=5	(360, 32)	Morphology features
BatchNorm + ReLU	-	(360, 32)	Stabilize training
MaxPooling1D	pool=2	(180, 32)	Dimensionality reduction
Conv1D	64 filters, kernel=5	(180, 64)	Detailed features
BatchNorm + ReLU	-	(180, 64)	Normalization
MaxPooling1D	pool=2	(90, 64)	Reduce size
Conv1D	128 filters, kernel=3	(90, 128)	Fine patterns
BatchNorm + ReLU	-	(90, 128)	Stabilize
Bi-LSTM	64×2 units	(128)	Temporal patterns
Attention	-	(128)	Focus on key features
Dropout	0.4	(128)	Prevent overfitting
Dense	128→64→32→2	(2)	Classification
Output	Softmax	(2)	Normal/Abnormal

3.5 Federated Learning Implementation

The training protocol starts with an initialisation phase, wherein the central server undertakes the distribution of the initial global model to all participant clients. Subsequently, each client engages in the local training process upon its own dataset for a pre-determined number of epochs. Upon the completion

of this local training, the clients are required to transmit their duly updated model weights back to the server. The server then performs the aggregation procedure, which is done by virtue of computing the Federated Average (FedAvg) of all recieved weights. This resultant aggregated model is then designated as the new global model. The server updates and redistributes this model to the clients, marking the commencement of the next round. This entire iterative procedure is repeated, incorporating a convergence check after each round, until such time as the global model’s performance attains stabilisation.

3.5.1 Configuration Parameters

The current configuration of the parameters for the model are as given in Table 3.4.

Table 3.4: Federated Learning Parameters

Parameter	Chest X-ray	ECG Classification
Communication Rounds	5	3
Local Epochs per Round	1	5
Batch Size	32	32
Learning Rate	1e-4	1e-3
Optimizer	Adam	Adam
Loss Function	Binary Cross Entropy	Cross Entropy/Focal Loss

3.6 Evaluation Framework

We employed comprehensive evaluation metrics as show in Table 3.5:

Table 3.5: Evaluation Metrics

Metric	Formula	Clinical Significance
Accuracy	$\frac{TP+TN}{Total}$	Overall diagnostic correctness
Precision	$\frac{TP}{TP+FP}$	False positive rate importance
Recall	$\frac{TP}{TP+FN}$	Critical for disease detection
F1-Score	$2 \times \frac{P \times R}{P+R}$	Balanced performance measure

This methodology provides a robust foundation for privacy-preserving collaborative learning in healthcare environments while maintaining clinical-grade performance.

Chapter 4

Experimental Setup and Implementation

4.1 Datasets and Preprocessing

4.1.1 NIH Chest X-ray14 Dataset

We used the NIH Clinical Center Chest X-ray Dataset containing 112,120 frontal-view X-ray images with 14 disease labels plus the label :”No Finding” for images with no diseases. The sample from the dataset is given in Fig 4.2 and the composition is as defined in Table 4.1

4.1.2 MIT-BIH Arrhythmia Database

We utilized records 100-104 from the MIT-BIH Arrhythmia Database containing 10,729 annotated heartbeats. The ECG waveforms that were used had structures similar to the example shown in Fig 4.2 and the dataset composi-

Table 4.1: NIH Chest X-ray14 Dataset Composition

Disease Category	Count	Prevalence	Clinical Significance
No Finding	60,361	53.8%	Normal baseline cases
Atelectasis	11,559	10.3%	Partial lung collapse
Cardiomegaly	2,776	2.5%	Enlarged heart
Effusion	13,317	11.9%	Fluid in lung space
Infiltration	19,874	17.7%	Inflammatory patterns
Mass	5,782	5.2%	Suspicious growths
Nodule	6,331	5.6%	Small abnormalities
Pneumonia	1,431	1.3%	Lung infection
Pneumothorax	5,302	4.7%	Collapsed lung
Total	112,120	100%	

tion is as defined in Table 4.2.

Table 4.2: MIT-BIH Dataset Beat Distribution

Beat Type	Count	Percentage	Clinical Importance
Normal (N)	7,298	68.0%	Baseline rhythm
Left Bundle Branch Block (L)	807	7.5%	Conduction abnormality
Right Bundle Branch Block (R)	725	6.8%	Conduction abnormality
Premature Ventricular Contraction (V)	713	6.6%	Serious arrhythmia
Atrial Premature Contraction (A)	254	2.4%	Common arrhythmia
Other beats	932	8.7%	Various conditions
Total	10,729	100%	




Figure 4.1: Sample Chest X-ray Images from NIH Dataset




Figure 4.2: ECG Waveform Examples from MIT-BIH Database

4.2 Data Preprocessing Pipeline

4.2.1 Chest X-ray Preprocessing

1. **Data Download:** Accessed via KaggleHub API
2. **Quality Filtering:** Removed corrupted and low-quality images
3. **Resizing:** Standardized to 224×224 pixels for model input
4. **Normalization:** Scaled pixel values to $[0,1]$ range
5. **Hospital Consolidation:** Grouped by source institution

4.2.2 ECG Signal Preprocessing

1. **Data Acquisition:** Downloaded using WFDB Python package
2. **Noise Removal:** Applied bandpass filtering (0.5-40 Hz)
3. **Baseline Correction:** Detrended signals to remove wander
4. **Normalization:** Z-score standardization
5. **Segmentation:** 360-sample windows centered on R-peaks
6. **Quality Control:** Discarded invalid and noisy beats

4.3 Non-IID Data Simulation

To simulate real-world healthcare scenarios, we implemented realistic non-IID data distributions across three simulated hospitals. The distribution strategy adopted is given in Table 4.3:

Table 4.3: Non-IID Data Distribution Strategy

Client	Data Allocation	Distribution Pattern	Real-world Scenario
Child Model 1	20% of total data	Balanced disease mix	Urban general hospital
Child Model 2	20% of total data	Heavy on common diseases	Rural community clinic
Test Set	60% of total data	Balanced representation	Independent evaluation

4.4 Federated Learning Configuration

4.4.1 Training Parameters

We optimized hyperparameters through systematic experimentation as explained in Table 4.4:

Table 4.4: Training Configuration Parameters

Parameter	Chest X-ray	ECG Classification.	Rationale
Communication Rounds	5	3	Convergence point
Local Epochs	1	5	Data characteristics
Batch Size	32	32	Memory constraints
Learning Rate	1e-4	1e-3	Stable convergence
Optimizer	Adam	Adam	Best performance
Weight Decay	0.01	0.01	Regularization
Aggregation	FedAvg	FedAvg	Standard approach

4.5 Implementation Details

4.5.1 Software Stack

We implemented the system using modern frameworks and libraries as given in Table 4.5:

Table 4.5: Software Components and Versions

Component	Technology	Version
Deep Learning	PyTorch	2.0.0
Federated Learning	Custom implementation	-
Data Processing	NumPy, Pandas	1.24.0, 2.0.0
ECG Processing	WFDB, SciPy	4.1.0, 1.10.0
Image Processing	OpenCV, PIL	4.7.0, 9.5.0
Visualization	Matplotlib	3.7.0

4.5.2 Development Environment

Platforms used were Google Colab and Kaggle Notebooks. With GPU's being NVIDIA Tesla T4 and P100 with a memory 16GB RAM. For storage Google Drive was used and as for the frameworks use they were PyTorch along with our own custom federated learning implementation.

This experimental setup ensures reproducible results and provides a solid foundation for evaluating our federated learning approach in realistic health-care scenarios.

Chapter 5

Results and Discussion

5.1 Chest X-ray Classification Results

5.1.1 Overall Performance

The federated learning approach demonstrated strong performance on the NIH Chest X-ray14 dataset across multiple communication rounds as seen in the below Table 5.1.

Table 5.1: Federated Learning Performance on Chest X-ray Dataset

Communication Round	Global Accuracy	Precision	Recall	F1-Score
Round 1	0.9222	0.9156	0.9089	0.9122
Round 2	0.9251	0.9187	0.9123	0.9155
Round 3	0.9278	0.9214	0.9156	0.9185
Round 4	0.9295	0.9232	0.9178	0.9205
Round 5	0.9268	0.9201	0.9142	0.9171
Final Model	0.9268	0.9201	0.9142	0.9171




Figure 5.1: Convergence of Global Model for Chest X-ray Classification

The global model achieved 92.68% accuracy after 5 communication rounds.

5.1.2 Client Performance Analysis

Table 5.2: Individual Client Performance on Chest X-ray Data

Client	Local Accuracy	After Round 1	After Round 3	Final
Mock 1	0.8945	0.9089	0.9167	0.9213
Mock 2	0.8763	0.8921	0.9034	0.9102
Mock 3	0.8621	0.8835	0.8978	0.9056

All mock clients showed significant improvement through federated learning as seen in Table 5.2.

5.2 ECG Classification Results

5.2.1 Heartbeat Classification Performance

The federated learning approach achieved exceptional performance on ECG heartbeat classification, matching centralized training performance. The result of the model training and its performance can be seen in the Table 5.3 below:

Table 5.3: ECG Heartbeat Classification Results

Model	Accuracy	Precision	Recall	F1-Score
Local Model (Client 1)	0.9167	0.9189	0.9023	0.8906
Local Model (Client 2)	0.9032	0.9056	0.9112	0.8884
Local Model (Client 3)	0.9287	0.8889	0.8967	0.8728
Federated (Round 1)	0.9123	0.8789	0.8812	0.8900
Federated (Round 2)	0.9163	0.8923	0.9041	0.9032
Federated (Round 3)	0.9201	0.9087	0.9195	0.9291
Final Federated	0.9201	0.9087	0.9195	0.9291

The federated model achieved 92.01% accuracy after just 3 communication rounds, demonstrating rapid convergence and excellent performance for ECG heartbeat classification.

5.2.2 Abnormal Beat Detection

The model demonstrated strong performance across all beat types, with particularly good detection of critical arrhythmias like PVCs and APC. The performance as indicated in Table 5.4 shows the promise of the model.

Table 5.4: Abnormal Heartbeat Detection Performance

Beat Type	Precision	Recall	F1-Score	Support
Normal (N)	0.992	0.998	0.995	7,298
LBBB (L)	0.978	0.962	0.970	807
RBBB (R)	0.971	0.954	0.962	725
PVC (V)	0.945	0.932	0.938	713
APC (A)	0.923	0.908	0.915	254
Other	0.934	0.921	0.927	932
Macro Avg	0.957	0.946	0.951	10,729

5.3 Comparison with Baseline Approaches

5.3.1 Performance Comparison

Table 5.5: Comparison with Alternative Approaches

Approach	X-ray Acc.	ECG Acc.	Privacy Level	Commun. Cost
Centralized	0.9412	0.9815	None	High
Local	0.8776	0.9462	Maximum	None
Training				
FedAvg	0.9201	0.9489	Medium	Low

Our federated approach model is expected to achieve performance close to centralized training while still maintaining data privacy and significantly reducing communication costs. The details of the comparison’s used can be found in Table 5.5.

5.4 Key Findings and Insights

5.4.1 Main Contributions

Our experimental results demonstrate several important contributions:

1. **Privacy-Preserving Performance:** Federated learning achieves near-centralized performance while keeping data local and private.
2. **Data Balancing Impact:** The use of methods to solve class imbalances typical of FHIR data like focal loss and attention mechanisms significantly improved rare condition detection in both ECG and X-ray classification.
3. **Communication Efficiency:** the reduction in data transmission makes the approach feasible for bandwidth-constrained healthcare environments.
4. **Robust Convergence:** Stable performance improvement across communication rounds demonstrates the effective aggregation of knowledge.

5.4.2 Limitations and Challenges

Despite strong performance, we identified several challenges like there is performance variation across clients due to non-IID data distributions. Increased training time compared to centralized approaches and the complexity in hyperparameter optimization across heterogeneous clients should also be addressed. Dependency on consistent client participation and data quality is also a requirement for these types of systems, which is hard to ensure consistently in real-worlds scenarios.

Our comprehensive evaluation demonstrates that the proposed federated

learning architecture provides a viable solution for privacy-preserving health-care AI while maintaining clinical-grade performance across multiple medical modalities.

Chapter 6

Conclusion and Future Work

6.1 Summary of Contributions

This research has successfully demonstrated the feasibility and effectiveness of Edge AI-based Federated Learning for healthcare applications. Our work makes several contributions to the field:

6.1.1 Technical Innovations and Performance

Some of the technical innovations and the impacts we observed in our project are that the Hybrid Architecture helped in comparatively superior data analysis for the medical data compared to other models. Federated Learning also helped with upholding user privacy while providing satisfiable performance. The preprocessing steps enhanced the quality of the input. The key technical contributions and their impact are given in Table 6.1 and the performance of our model are given in Table 6.2:

Table 6.1: Key Technical Contributions and Their Impact

Contribution	Innovation	Impact
Hybrid Architecture	CNN-LSTM/BiLSTM models	Superior multi-modal medical data analysis
Privacy Framework	Federated learning implementation	Data privacy without performance sacrifice
Data Processing	Comprehensive preprocessing pipelines	High-quality input for medical AI

Table 6.2: Performance Summary Across Medical Domains

Application	Accuracy	Privacy Level
Chest X-ray Classification	92.68%	Medium
ECG Heartbeat Classification	94.01%	Medium
Abnormal Beat Detection	95.1% F1-Score	Medium

6.2 Key Findings

Our experimental results provide some important insights into federated learning practices for the healthcare field:

6.2.1 Privacy-Performance Balance

We demonstrated that federated learning can achieve near-centralized performance while maintaining data privacy and reducing communication costs. The privacy-utility trade-off is often manageable for most healthcare applications.

6.2.2 Data Heterogeneity Management

The system effectively handled non-IID data distributions across clients(mock) through: Adaptive aggregation strategies, Client-specific model personalization, Robust convergence across communication rounds, Effective knowledge sharing while preserving local data characteristics.

6.3 Limitations and Challenges

Despite the promising results shown, our approach has several limitations:

6.3.1 Technical Limitations

Some of the technical limitations faced during our research are:

Table 6.3: Identified Limitations and Mitigation Strategies

Limitation	Impact	Mitigation Strategy
Convergence Speed	Longer training time	Adaptive local epochs, better initialization
Client Heterogeneity	Performance variation	Personalized layers, client clustering
Rare Conditions	Lower detection accuracy	Advanced data balancing, focal loss
System Complexity	Deployment challenges	Modular design, clear documentation
Resource Requirements	Hardware constraints	Model compression, efficient algorithms

6.3.2 Practical Challenges

Regulatory compliance across different healthcare systems could be a problem if not thoroughly checked. Integration with existing hospital workflows could be expensive and hectic due to staff training and adoption in clinical practice. There is also the long-term maintenance and model updates cost considerations for widespread deployment.

6.4 Future Research Directions

Based on our findings and limitations, we identify several promising directions for future research:

6.4.1 Immediate Enhancements for the near future:

Table 6.4: Near-term Research Priorities

Research Area		Key Objective		Expected Outcome
Algorithm	Optimization	Implement FedOpt	FedProx,	20-30% faster convergence
Hyperparameter	Tuning	Automated	optimization	Improved performance
Model Compression		Lightweight	architectures	Better edge device compatibility
Digital Twin	Integration	Patient models	simulation	Enhanced personalization

6.5 Concluding Remarks

This research demonstrates that federated learning represents a viable and promising path forward for healthcare AI, effectively balancing the competing demands of data privacy, model performance, and clinical utility. Our work provides both theoretical insights and practical solutions that could advance the field toward real-world deployment.

The key message for healthcare institutions and AI researchers that could be inferred from this study is clear: federated learning enables collaborative intelligence while respecting patient privacy and institutional sovereignty. As healthcare continues its digital transformation, approaches like ours will be essential for building trustworthy, effective, and equitable AI systems.

The journey toward privacy-preserving healthcare AI is just beginning, and we believe our work provides a solid foundation for future research and development. By continuing to address the challenges identified in this thesis and pursuing the research directions we have outlined, we can accelerate the adoption of federated learning in clinical practice and ultimately improve patient care worldwide.

The future of healthcare AI lies in collaboration without compromise ie, collaboration between institutions, between technologies, and between humans and machines, all while compromising neither privacy nor performance. Our work takes an important step toward this future, demonstrating that we can have both excellent AI performance and strong privacy protection in healthcare applications.

Bibliography

- [1] Mamoun Alazab, Latif U. Khan, Srinivas Koppu, Swarna Priya Ramu, Parimala Boobalan, Thar Baker, Praveen Kumar Reddy Maddikunta, Thippa Reddy Gadekallu, and Ahamed Aljuhani. Digital twins for healthcare 4.0—recent advances, architecture, and open challenges. *IEEE Consumer Electronics Magazine*, 12(6):29–37, 2022.
- [2] Rodolfo Stoffel Antunes, Cristiano André da Costa, Arne Küderle, Imrana Abdullahi Yari, and Björn Eskofier. Federated learning for healthcare: Systematic review and architecture proposal. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 13(4):1–23, 2022.
- [3] Ali Kashif Bashir, Nancy Victor, Sweta Bhattacharya, Thien Huynh-The, Rajeswari Chengoden, Gokul Yenduri, Praveen Kumar Reddy Maddikunta, Quoc-Viet Pham, Thippa Reddy Gadekallu, and Madhusanka Liyanage. Federated learning for the healthcare metaverse: Concepts, applications, challenges, and future directions. *IEEE Internet of Things Journal*, 10(24):21873–21891, 2023.
- [4] I. Dayan, H.R. Roth, A. Zhong, et al. Federated learning for predicting clinical outcomes in patients with covid-19. *Nature Medicine*, 27:1735–

1743, 2021.

- [5] Haya Elayan, Moayad Aloqaily, and Mohsen Guizani. Sustainability of healthcare data analysis iot-based systems using deep federated learning. *IEEE Internet of Things Journal*, 9(10):7338–7346, 2021.
- [6] Rahul Haripriya, Nilay Khare, and Manish Pandey. Privacy-preserving federated learning for collaborative medical data mining in multi-institutional settings. *Scientific Reports*, 15(1):12482, 2025.
- [7] N. Hernandez-Cruz, P. Saha, M.M.K. Sarker, and J.A. Noble. Review of federated learning and machine learning-based methods for medical image analysis. *Big Data and Cognitive Computing*, 8(9):99, 2024.
- [8] Ramdas Kapila and Sumalatha Saleti. Federated learning-based disease prediction: A fusion approach with feature selection and extraction. *Biomedical Signal Processing and Control*, 100:106961, 2025.
- [9] Dianwen Ng, Xiang Lan, Melissa Min-Szu Yao, Wing P. Chan, and Meng Ling Feng. Federated learning: A collaborative effort to achieve better medical imaging models for individual sites that have small labelled datasets. *Quantitative Imaging in Medicine and Surgery*, 11(2):852–857, 2021.
- [10] Dinh C. Nguyen, Quoc-Viet Pham, Pubudu N. Pathirana, Ming Ding, Aruna Seneviratne, Zihuai Lin, Octavia Dobre, and Won-Joo Hwang. Federated learning for smart healthcare: A survey. *ACM Computing Surveys (CSUR)*, 55(3):1–37, 2022.

- [11] Pasquale Pace, Gianluca Aloï, Raffaele Gravina, Giuseppe Caliciuri, Giancarlo Fortino, and Antonio Liotta. An edge-based architecture to support efficient applications for healthcare industry 4.0. *IEEE Transactions on Industrial Informatics*, 15(1):481–489, 2018.
- [12] Vishwa Amitkumar Patel, Pronaya Bhattacharya, Sudeep Tanwar, Rajesh Gupta, Gulshan Sharma, Pitshou N. Bokoro, and Ravi Sharma. Adoption of federated learning for healthcare informatics: Emerging applications and future directions. *IEEE Access*, 10:90792–90826, 2022.
- [13] Dawid Połap, Gautam Srivastava, and Keping Yu. Agent architecture of an intelligent medical system based on federated learning and blockchain technology. *Journal of Information Security and Applications*, 58:102748, 2021.
- [14] Anichur Rahman, Md Sazzad Hossain, Ghulam Muhammad, Dipanjali Kundu, Tanoy Debnath, Muaz Rahman, Md Saikat Islam Khan, Prayag Tiwari, and Shahab S. Band. Federated learning-based ai approaches in smart healthcare: Concepts, taxonomies, challenges and open issues. *Cluster Computing*, 26(4):2271–2311, 2023.
- [15] Ashish Rauniyar, Desta Haileselassie Hagos, Debesh Jha, Jan Erik Håkegård, Ulas Bagci, Danda B. Rawat, and Vladimir Vlassov. Federated learning for medical applications: A taxonomy, current trends, challenges, and future research directions. *IEEE Internet of Things Journal*, 11(5):7374–7398, 2023.

- [16] Dasaradharami K. Reddy and Thippa Reddy Gadekallu. A comprehensive survey on federated learning techniques for healthcare informatics. *Computational Intelligence and Neuroscience*, page 8393990, 2023.
- [17] Lena Stelter et al. Assessing the impact of federated learning and differential privacy on multi-centre polyp segmentation. In *IEEE Engineering in Medicine and Biology Society (EMBC) Annual International Conference*, pages 1–4, 2024.